

Assignment 3

```
[cloudera@quickstart DSBDA]$ hive
```

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
```

```
--CREATING AND USING SPECIFIC DATABASE
```

```
hive> create database flightDB;
```

```
OK
```

```
Time taken: 0.329 seconds
```

```
hive> use flightDB;
```

```
OK
```

```
Time taken: 0.02 seconds
```

```
--CREATING AND ALTERING TABLE
```

```
hive> create table flight_info(
```

```
> flight_id string,
```

```
> origin string,
```

```
> destination string,
```

```
> delay int
```

```
> )
```

```
> ROW FORMAT DELIMITED
```

```
> FIELDS TERMINATED BY ','
```

```
> LINES TERMINATED BY '\n'
```

```
> STORED AS TEXTFILE;|
```

```
OK
```

```
Time taken: 0.511 seconds
```

```
hive> alter table flight_info change column delay total_delay int;
```

```
OK
```

```
Time taken: 0.261 seconds
```

```
hive> desc flight_info;
```

```
OK
```

```
flight_id          string
```

```
origin             string
```

```
destination        string
```

```
total_delay        int
```

```
Time taken: 0.129 seconds, Fetched: 4 row(s)
```

```
hive> select * from flight_info;
```

```
OK
```

```
Time taken: 0.455 seconds
```

```
hive> alter table flight_info add columns(
```

```
> departure_date date
```

```
> );
```

```
OK
```

```
Time taken: 0.122 seconds
```

```
hive> desc flight_info;
```

```
OK
```

```
flight_id          string
```

```
origin             string
```

```
destination        string
```

```
total_delay        int
```

```
departure_date     date
```

```
Time taken: 0.085 seconds, Fetched: 5 row(s)
```

```
hive> select * from flight_info;
```

```
OK
```

```
Time taken: 0.058 seconds
```

```
--CREATING AN EXTERNAL HIVE TABLE
```

```
hive> create external table flight_info2(
```

```
> flight_id int,
```

```
> origin string,
```

```
> des string,
```

```
> delay int
```

```
> )
```

```
> ROW FORMAT DELIMITED
```

```
> FIELDS TERMINATED BY ','
```

```
> LOCATION '/home/cloudera/Desktop/DSBDA';
```

```
OK
```

```
Time taken: 0.11 seconds
```

```
hive> select * from flight_info;
```

```
OK
```

```
Time taken: 0.057 seconds
```

```
--LOADING DATA FROM EXTERNAL FILE INTO THE TABLE
```

```
hive> load data local inpath '/home/cloudera/Desktop/DSBDA/test.csv' overwrite into table flight_info;
```

```
Loading data to table flightdb.flight_info
```

```
Table flightdb.flight_info stats: [numFiles=1, numRows=0, totalSize=279, rawDataSize=0]
```

```
OK
```

```
Time taken: 0.667 seconds
```

```
--CHECKING IF WORKED
```

```
hive> select * from flight_info;
```

```
OK
```

```
F001  Mumabi  Delhi  45      2025-02-15
```

```
F002  Mumabi  Nagpur  -15     2025-02-17
```

```
F003  Mumabi  Indore  30      2025-02-18
```

```
F004  Kolkata  Inodore  0       2025-02-20
```

```
F005  Kochin   Srinagar  60      2025-02-22
```

F006	Hyderabad	Lucknow	0	2025-02-25
F007	Mumabi	Ahemdabad	20	2025-02-26
F008	Delhi	Ahemdabad	-20	2025-02-26
	NULL	NULL	NULL	NULL
	NULL	NULL	NULL	NULL
	NULL	NULL	NULL	NULL

Time taken: 0.056 seconds, Fetched: 11 row(s)

--INSERTING DATA MANUALLY INTO THE TABLE

hive> insert into table flight_info values

> ('F009','Mumbai','Pune',0,'2025-02-28');

Query ID = cloudera_20250205205656_91181bc3-2937-4856-95f3-dce6200c6044

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1738815374210_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/

application_1738815374210_0001/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1738815374210_0001

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2025-02-05 20:56:24,686 Stage-1 map = 0%, reduce = 0%

2025-02-05 20:56:31,452 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.07 sec

MapReduce Total cumulative CPU time: 1 seconds 70 msec

Ended Job = job_1738815374210_0001

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/flightdb.db/flight_info/.hive-

staging_hive_2025-02-05_20-56-11_634_8364896290987674265-1/-ext-10000

Loading data to table flightdb.flight_info

Table flightdb.flight_info stats: [numFiles=2, numRows=1, totalSize=309, rawDataSize=29]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 1.07 sec HDFS Read: 4906 HDFS Write: 106 SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 70 msec

OK

Time taken: 21.308 seconds

hive> select * from flight_info;

OK

F009	Mumbai	Pune	0	2025-02-28
F001	Mumabi	Delhi	45	2025-02-15
F002	Mumabi	Nagpur	-15	2025-02-17
F003	Mumabi	Indore	30	2025-02-18
F004	Kolkata	Inodore	0	2025-02-20
F005	Kochin	Srinagar	60	2025-02-22
F006	Hyderabad	Lucknow	0	2025-02-25
F007	Mumabi	Ahemdabad	20	2025-02-26
F008	Delhi	Ahemdabad	-20	2025-02-26
	NULL	NULL	NULL	NULL
	NULL	NULL	NULL	NULL
	NULL	NULL	NULL	NULL

Time taken: 0.056 seconds, Fetched: 12 row(s)

--CREATING INDEX ON COLUMN 'departure_date' IN TABLE

hive> create index flight_index

> on table flight_info(departure_date)

> as 'compact'

> with deferred rebuild;

OK

Time taken: 0.291 seconds

--SHOWING THE INDEX

hive> show index on flight_info;

OK

flight_index	flight_info	departure_date	flightdb__flight_info_flight_index__
compact			

Time taken: 0.083 seconds, Fetched: 1 row(s)

hive> describe formatted flight_info;

OK

#	col_name	data_type	comment
	flight_id	string	
	origin	string	
	destination	string	
	total_delay	int	
	departure_date	date	

Detailed Table Information

Database: flightdb

Owner: cloudera

CreateTime: Wed Feb 05 20:38:48 PST 2025

LastAccessTime: UNKNOWN

Protect Mode: None

Retention: 0

Location: hdfs://quickstart.cloudera:8020/user/hive/warehouse/flightdb.db/flight_info

Table Type: MANAGED_TABLE

Table Parameters:

COLUMN_STATS_ACCURATE true

last_modified_by cloudera

```

        last_modified_time      1738817192
        numFiles                2
        numRows                 1
        rawDataSize              29
        totalSize                309
        transient_lastDdlTime    1738817792

# Storage Information
SerDe Library:      org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:        org.apache.hadoop.mapred.TextInputFormat
OutputFormat:        org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:          No
Num Buckets:         -1
Bucket Columns:      []
Sort Columns:         []
Storage Desc Params:
    field.delim       ,
    line.delim         \n
    serialization.format
Time taken: 0.169 seconds, Fetched: 39 row(s)

--CALCULATING AVERAGE DELAY IN MONTH OF FEB
hive> select avg(delay) as avg_delay_feb
    > from flight_info
    > ;
FAILED: SemanticException [Error 10004]: Line 1:11 Invalid table alias or column reference 'delay': (possible
column names are: flight_id, origin, destination, total_delay, departure_date)
hive> select avg(total_delay) as avg_delay_feb
    > from flight_info;
Query ID = cloudera_20250205210404_b188b977-2bed-4ce7-a00e-4acc46f4632f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
    set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
    set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
    set mapreduce.job.reduces=<number>
Starting Job = job_1738815374210_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/
application_1738815374210_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1738815374210_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2025-02-05 21:04:31,441 Stage-1 map = 0%,  reduce = 0%
2025-02-05 21:04:37,792 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 0.89 sec
2025-02-05 21:04:45,214 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 1.8 sec
MapReduce Total cumulative CPU time: 1 seconds 800 msec
Ended Job = job_1738815374210_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.8 sec HDFS Read: 9025 HDFS Write: 19 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 800 msec
OK
13.333333333333334
Time taken: 22.405 seconds, Fetched: 1 row(s)
hive> exit;

```