

Operations:

- Data cleaning
- Data integration
- Data transformation
- Error correcting
- Data Model Backend

```
import pandas as pd
import numpy as np

df = pd.read_csv('airquality_data.csv', encoding='cp1252')

C:\Users\Gayatri Tagalpallewa\AppData\Local\Temp\
ipykernel_21076\2182913842.py:1: DtypeWarning: Columns (0) have mixed
types. Specify dtype option on import or set low_memory=False.
  df = pd.read_csv('airquality_data.csv', encoding='cp1252')

df.head()
```

	stn_code	sampling_date	state	location	agency	\
0	150.0	February - M021990	Andhra Pradesh	Hyderabad	NaN	
1	151.0	February - M021990	Andhra Pradesh	Hyderabad	NaN	
2	152.0	February - M021990	Andhra Pradesh	Hyderabad	NaN	
3	150.0	March - M031990	Andhra Pradesh	Hyderabad	NaN	
4	151.0	March - M031990	Andhra Pradesh	Hyderabad	NaN	

	type	so2	no2	rspm	spm	\
0	Residential, Rural and other Areas	4.8	17.4	NaN	NaN	
1	Industrial Area	3.1	7.0	NaN	NaN	
2	Residential, Rural and other Areas	6.2	28.5	NaN	NaN	
3	Residential, Rural and other Areas	6.3	14.7	NaN	NaN	
4	Industrial Area	4.7	7.5	NaN	NaN	

	location_monitoring_station	pm2_5	date
0	NaN	NaN	1990-02-01
1	NaN	NaN	1990-02-01
2	NaN	NaN	1990-02-01
3	NaN	NaN	1990-03-01
4	NaN	NaN	1990-03-01

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 435742 entries, 0 to 435741
```

```
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   stn_code                             291665 non-null object
1   sampling_date                         435739 non-null object
2   state                                435742 non-null object
3   location                             435739 non-null object
4   agency                               286261 non-null object
5   type                                 430349 non-null object
6   so2                                  401096 non-null float64
7   no2                                  419509 non-null float64
8   rspm                                 395520 non-null float64
9   spm                                  198355 non-null float64
10  location_monitoring_station           408251 non-null object
11  pm2_5                                9314 non-null float64
12  date                                  435735 non-null object
dtypes: float64(5), object(8)
memory usage: 43.2+ MB

df.columns

Index(['stn_code', 'sampling_date', 'state', 'location', 'agency',
      'type',
      'so2', 'no2', 'rspm', 'spm', 'location_monitoring_station',
      'pm2_5',
      'date'],
      dtype='object')
```

Data Cleaning

```
# Change data type from float64 to float32 for Space Complexity
df['so2'] = df['so2'].astype('float32')
df['no2'] = df['no2'].astype('float32')
df['rspm'] = df['rspm'].astype('float32')
df['spm'] = df['spm'].astype('float32')
df['date'] = df['date'].astype('string')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 435742 entries, 0 to 435741
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	stn_code	291665 non-null	object
1	sampling_date	435739 non-null	object
2	state	435742 non-null	object
3	location	435739 non-null	object
4	agency	286261 non-null	object
5	type	430349 non-null	object

```
6    so2                401096 non-null float32
7    no2                419509 non-null float32
8    rspm              395520 non-null float32
9    spm               198355 non-null float32
10   location_monitoring_station 408251 non-null object
11   pm2_5             9314 non-null float64
12   date              435735 non-null string
dtypes: float32(4), float64(1), object(7), string(1)
memory usage: 36.6+ MB
```

```
df=df.drop_duplicates()
```

```
df.isna().sum()
```

```
stn_code                144077
sampling_date              3
state                    0
location                 3
agency                 149466
type                    5357
so2                     34632
no2                     16222
rspm                    40035
spm                    236908
location_monitoring_station 27303
pm2_5                   425754
date                      7
dtype: int64
```

```
percent_missing = df.isnull().sum() * 100 / len(df)
```

```
percent_missing.sort_values(ascending=False)
```

```
pm2_5                97.859185
spm                 54.453097
agency              34.354630
stn_code            33.115973
rspm                9.202010
so2                 7.960135
location_monitoring_station 6.275571
no2                 3.728613
type                1.231302
date                0.001609
sampling_date       0.000690
location            0.000690
state               0.000000
dtype: float64
```

```
df=df.drop(['stn_code',
'agency', 'sampling_date', 'location_monitoring_station', 'pm2_5'], axis
= 1)
```

```
df.head()
```

	state	location	type	so2
no2 \				
0	Andhra Pradesh	Hyderabad	Residential, Rural and other Areas	4.8
17.4				
1	Andhra Pradesh	Hyderabad	Industrial Area	3.1
7.0				
2	Andhra Pradesh	Hyderabad	Residential, Rural and other Areas	6.2
28.5				
3	Andhra Pradesh	Hyderabad	Residential, Rural and other Areas	6.3
14.7				
4	Andhra Pradesh	Hyderabad	Industrial Area	4.7
7.5				

	rspm	spm	date
0	NaN	NaN	1990-02-01
1	NaN	NaN	1990-02-01
2	NaN	NaN	1990-02-01
3	NaN	NaN	1990-03-01
4	NaN	NaN	1990-03-01

```
df.columns
```

```
Index(['state', 'location', 'type', 'so2', 'no2', 'rspm', 'spm',  
      'date'], dtype='object')
```

```
col_var = ['state', 'location', 'type', 'date']  
col_num = ['so2', 'no2', 'rspm', 'spm']
```

```
for col in df.columns:  
    if df[col].dtype == 'object' or df[col].dtype == 'string':  
        df[col] = df[col].fillna(df[col].mode()[0])  
    else:  
        df[col] = df[col].fillna(df[col].mean())
```

```
df.isna().sum()
```

state	0
location	0
type	0
so2	0
no2	0
rspm	0
spm	0
date	0

```
dtype: int64  
  
df
```

	state	location \
0	Andhra Pradesh	Hyderabad
1	Andhra Pradesh	Hyderabad
2	Andhra Pradesh	Hyderabad
3	Andhra Pradesh	Hyderabad
4	Andhra Pradesh	Hyderabad
...
435737	West Bengal	ULUBERIA
435738	West Bengal	ULUBERIA
435739	andaman-and-nicobar-islands	Guwahati
435740	Lakshadweep	Guwahati
435741	Tripura	Guwahati

rspm \	type	so2	no2
0	Residential, Rural and other Areas	4.800000	17.400000
108.871712			
1	Industrial Area	3.100000	7.000000
108.871712			
2	Residential, Rural and other Areas	6.200000	28.500000
108.871712			
3	Residential, Rural and other Areas	6.300000	14.700000
108.871712			
4	Industrial Area	4.700000	7.500000
108.871712			
...
...			
435737	RIRU0	22.000000	50.000000
143.000000			
435738	RIRU0	20.000000	46.000000
171.000000			
435739	Residential, Rural and other Areas	10.830467	25.823299
108.871712			
435740	Residential, Rural and other Areas	10.830467	25.823299
108.871712			
435741	Residential, Rural and other Areas	10.830467	25.823299
108.871712			

	spm	date
0	220.774796	1990-02-01
1	220.774796	1990-02-01
2	220.774796	1990-02-01
3	220.774796	1990-03-01
4	220.774796	1990-03-01
...
435737	220.774796	2015-12-24
435738	220.774796	2015-12-29
435739	220.774796	2015-03-19
435740	220.774796	2015-03-19
435741	220.774796	2015-03-19

```
[435068 rows x 8 columns]
```

```
df.isna().sum()
```

```
state      0
location   0
type        0
so2         0
no2         0
rspm        0
spm         0
date        0
dtype: int64
```

Data integration

```
subSet1 = df[['state', 'type']]
subSet2 = df[['state', 'location']]
```

```
subSet1.head()
```

	state	type
0	Andhra Pradesh	Residential, Rural and other Areas
1	Andhra Pradesh	Industrial Area
2	Andhra Pradesh	Residential, Rural and other Areas
3	Andhra Pradesh	Residential, Rural and other Areas
4	Andhra Pradesh	Industrial Area

```
subSet2.head()
```

	state	location
0	Andhra Pradesh	Hyderabad
1	Andhra Pradesh	Hyderabad
2	Andhra Pradesh	Hyderabad
3	Andhra Pradesh	Hyderabad
4	Andhra Pradesh	Hyderabad

```
concatenated_df = pd.concat([subSet1, subSet2], axis=1)
```

```
concatenated_df
```

	state	type \
0	Andhra Pradesh	Residential, Rural and other Areas
1	Andhra Pradesh	Industrial Area
2	Andhra Pradesh	Residential, Rural and other Areas

```

3          Andhra Pradesh  Residential, Rural and other
Areas
4          Andhra Pradesh          Industrial
Area
...          ...          ..
.
435737      West Bengal
RIRU0
435738      West Bengal
RIRU0
435739  andaman-and-nicobar-islands  Residential, Rural and other
Areas
435740      Lakshadweep  Residential, Rural and other
Areas
435741      Tripura  Residential, Rural and other
Areas

          state  location
0      Andhra Pradesh  Hyderabad
1      Andhra Pradesh  Hyderabad
2      Andhra Pradesh  Hyderabad
3      Andhra Pradesh  Hyderabad
4      Andhra Pradesh  Hyderabad
...          ...          ...
435737      West Bengal  ULUBERIA
435738      West Bengal  ULUBERIA
435739  andaman-and-nicobar-islands  Guwahati
435740      Lakshadweep  Guwahati
435741      Tripura  Guwahati

[435068 rows x 4 columns]

```

Error Correcting

```

def remove_outliers(column):
    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)
    IQR = Q3 - Q1
    threshold = 1.5 * IQR
    outlier_mask = (column < Q1 - threshold) | (column > Q3 +
threshold)
    return column[~outlier_mask]

df.columns

Index(['state', 'location', 'type', 'so2', 'no2', 'rspm', 'spm',
'date'], dtype='object')

# Remove outliers for each column using a loop
col_name = ['so2', 'no2', 'rspm', 'spm']

```

```


for col in col_name:
    df[col] = remove_outliers(df[col])



import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(10, 6)) # Adjust the figure size if needed

for col in col_name:
    sns.boxplot(data=df[col])
    plt.title(col)
    plt.show()

```







Data Transform

```
from sklearn.preprocessing import LabelEncoder

col_label= ['state','location','type']
# Initialize LabelEncoder

encoder = LabelEncoder()
# Iterate over columns
for col in df.columns:
    # Fit and transform the column
    df[col] = encoder.fit_transform(df[col])
```

df

	state	location	type	so2	no2	rspm	spm	date
0	0	114	6	446	1489	2030	464	213
1	0	114	1	197	250	2030	464	213
2	0	114	6	790	3096	2030	464	213
3	0	114	6	823	1144	2030	464	214
4	0	114	1	427	301	2030	464	214
...
435737	35	282	3	2888	5307	2534	464	5059
435738	35	282	3	2809	5113	3098	464	5064
435739	36	100	6	1638	2696	2030	464	4779
435740	17	100	6	1638	2696	2030	464	4779

435741	31	100	6	1638	2696	2030	464	4779
--------	----	-----	---	------	------	------	-----	------

[435068 rows x 8 columns]

```
# import pandas library
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.linear_model import LogisticRegression
import seaborn as sns
import matplotlib.pyplot as plt

# Reading csv file
df = pd.read_csv("Heart.csv")
df.head()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6

	caa	thall	output
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

Data Cleaning

```
df = df.drop_duplicates()

# Count ,min,max ,etc of each column
df.describe()
```

	age	sex	cp	trtbps	chol
count	302.000000	302.000000	302.000000	302.000000	302.000000
mean	54.42053	0.682119	0.963576	131.602649	246.500000
std	9.04797	0.466426	1.032044	17.563394	51.753489

```

0.356686
min      29.000000      0.000000      0.000000      94.000000      126.000000
0.000000
25%      48.000000      0.000000      0.000000      120.000000      211.000000
0.000000
50%      55.500000      1.000000      1.000000      130.000000      240.500000
0.000000
75%      61.000000      1.000000      2.000000      140.000000      274.750000
0.000000
max      77.000000      1.000000      3.000000      200.000000      564.000000
1.000000

```

```

          restecg      thalachh          exng      oldpeak          slp
caa \
count  302.000000  302.000000  302.000000  302.000000  302.000000
302.000000
mean    0.526490  149.569536    0.327815    1.043046    1.397351
0.718543
std     0.526027   22.903527    0.470196    1.161452    0.616274
1.006748
min     0.000000   71.000000    0.000000    0.000000    0.000000
0.000000
25%     0.000000  133.250000    0.000000    0.000000    1.000000
0.000000
50%     1.000000  152.500000    0.000000    0.800000    1.000000
0.000000
75%     1.000000  166.000000    1.000000    1.600000    2.000000
1.000000
max     2.000000  202.000000    1.000000    6.200000    2.000000
4.000000

```

```

          thall          output
count  302.000000  302.000000
mean    2.314570    0.543046
std     0.613026    0.498970
min     0.000000    0.000000
25%     2.000000    0.000000
50%     2.000000    1.000000
75%     3.000000    1.000000
max     3.000000    1.000000

```

```

# Information about each column data
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 302 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         302 non-null   int64

```

1	sex	302	non-null	int64
2	cp	302	non-null	int64
3	trtbps	302	non-null	int64
4	chol	302	non-null	int64
5	fbs	302	non-null	int64
6	restecg	302	non-null	int64
7	thalachh	302	non-null	int64
8	exng	302	non-null	int64
9	oldpeak	302	non-null	float64
10	slp	302	non-null	int64
11	caa	302	non-null	int64
12	thall	302	non-null	int64
13	output	302	non-null	int64

dtypes: float64(1), int64(13)

memory usage: 35.4 KB

#Finding null values in each column

df.isna().sum()

age	0
sex	0
cp	0
trtbps	0
chol	0
fbs	0
restecg	0
thalachh	0
exng	0
oldpeak	0
slp	0
caa	0
thall	0
output	0

dtype: int64

Data Integration

df.head()

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6

2

	caa	thall	output
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
df.fbs.unique()
```

```
array([1, 0], dtype=int64)
```

```
subSet1 = df[['age', 'cp', 'chol', 'thalachh']]
```

```
subSet2 = df[['exng', 'slp', 'output']]
```

```
merged_df = subSet1.merge(right=subSet2, how='cross')  
merged_df.head()
```

	age	cp	chol	thalachh	exng	slp	output
0	63	3	233	150	0	0	1
1	63	3	233	150	0	0	1
2	63	3	233	150	0	2	1
3	63	3	233	150	0	2	1
4	63	3	233	150	1	2	1

Error Correcting


```
df.columns
```



```
Index(['age', 'sex', 'cp', 'trtbps', 'chol', 'fbs', 'restecg',  
      'thalachh',  
      'exng', 'oldpeak', 'slp', 'caa', 'thall', 'output'],  
      dtype='object')
```


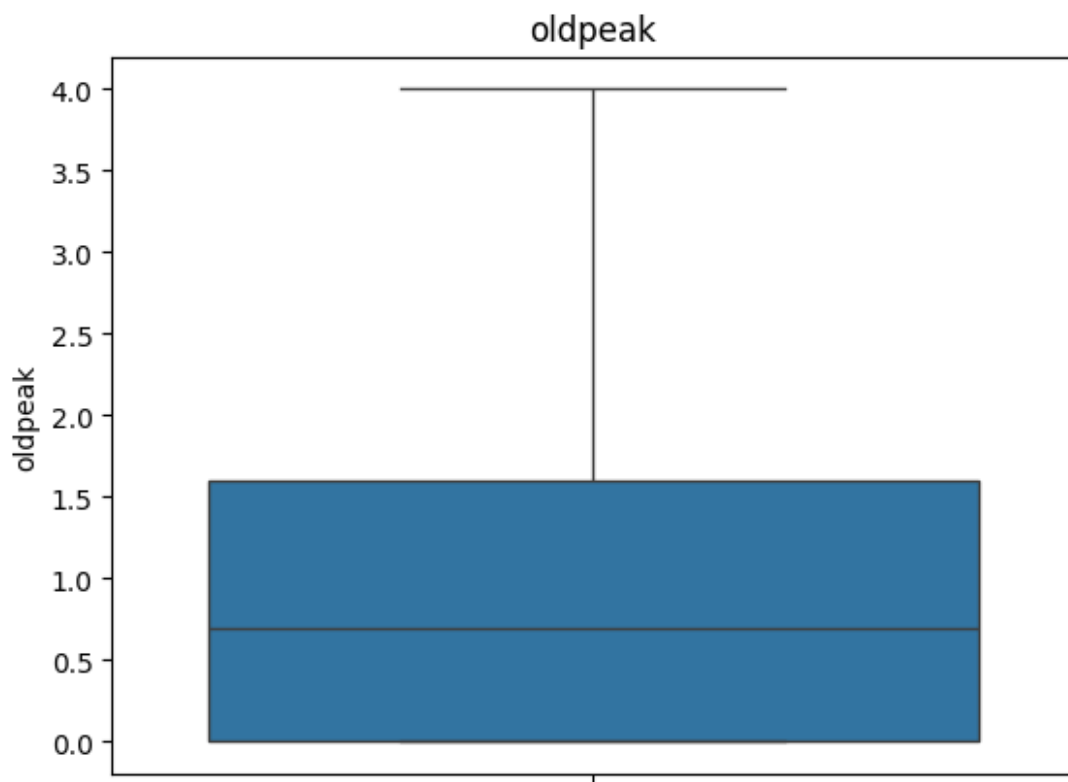
```
def remove_outliers(column):  
    Q1 = column.quantile(0.25)  
    Q3 = column.quantile(0.75)  
    IQR = Q3 - Q1  
    threshold = 1.5 * IQR  
    outlier_mask = (column < Q1 - threshold) | (column > Q3 +  
threshold)  
    return column[~outlier_mask]
```


```
# Remove outliers for each column using a loop  
col_name = ['cp', 'thalachh', 'exng', 'oldpeak', 'slp', 'caa']  
for col in col_name:  
    df[col] = remove_outliers(df[col])
```

```
plt.figure(figsize=(10, 6)) # Adjust the figure size if needed
for col in col_name:
    sns.boxplot(data=df[col])
    plt.title(col)
    plt.show()
```









```
df = df.dropna()
df.isna().sum()
age          0
sex          0
cp           0
trtbps      0
chol        0
fbs         0
restecg     0
thalachh    0
exng        0
oldpeak     0
slp         0
caa         0
thall       0
output      0
dtype: int64

df = df.drop('fbs',axis=1)

# Compute correlations between features and target
correlations = df.corr()['output'].drop('output')

# Print correlations
```

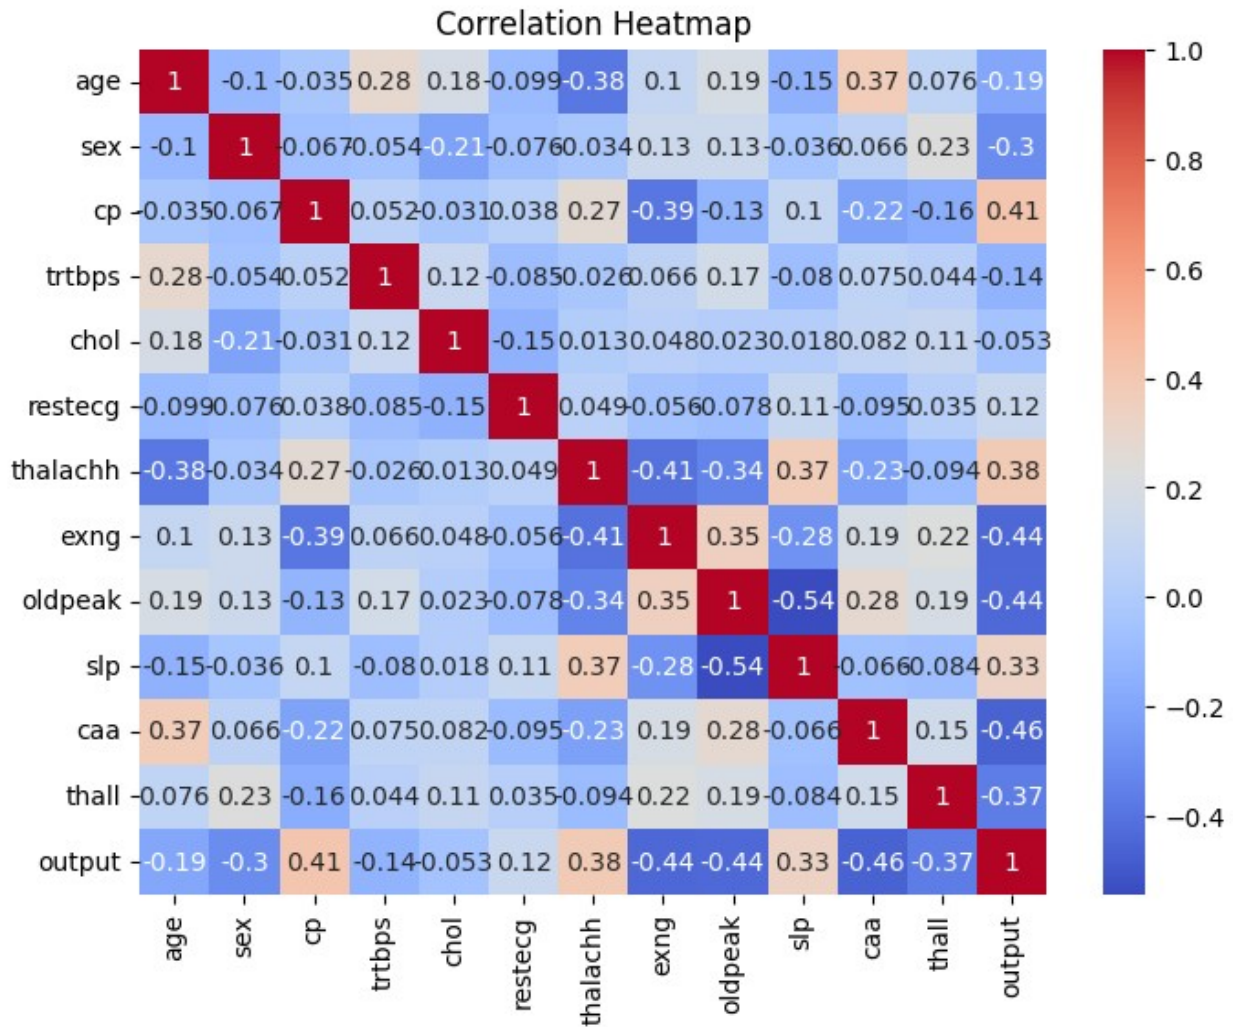
```
print("Correlation with the Target:")
print(correlations)
print()

# Plot correlation heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```

Correlation with the Target:

age	-0.193798
sex	-0.303271
cp	0.410807
trtbps	-0.135238
chol	-0.052796
restecg	0.122071
thalachh	0.384609
exng	-0.444401
oldpeak	-0.437895
slp	0.329432
caa	-0.460816
thall	-0.366390

Name: output, dtype: float64



```
# df.isna().sum()
```

Data Split

```
# splitting data using train test split
x = df[['cp', 'thalachh', 'exng', 'oldpeak', 'slp', 'caa']]
y = df.output
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)

x_train.shape, x_test.shape, y_train.shape, y_test.shape
((220, 6), (55, 6), (220,), (55,))
```

Data transformation

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()

x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
```

Data model building

```
y_train= np.array(y_train).reshape(-1, 1)
y_test= np.array(y_test).reshape(-1, 1)
```

```
y_train.shape
```

```
(220, 1)
```

```
model = LogisticRegression()
model.fit(x_train_scaled, y_train)
```

```
# Make predictions on the test set
y_pred = model.predict(x_test_scaled)
```

```
# Evaluate the model's accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

```
Accuracy: 0.8363636363636363
```

```
C:\Users\Gayatri Tagalpallewa\AppData\Roaming\Python\Python312\site-
packages\sklearn\utils\validation.py:1339: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change
the shape of y to (n_samples, ), for example using ravel().
```

```
    y = column_or_1d(y, warn=True)
```

```
#Classification model using Decision Tree
```

```
from sklearn.tree import DecisionTreeClassifier
tc=DecisionTreeClassifier(criterion='entropy')
tc.fit(x_train_scaled,y_train)
y_pred=tc.predict(x_test_scaled)
```

```
print("Training Accuracy Score :",accuracy_score(y_pred,y_test))
print("Training Confusion Matrix  :",confusion_matrix(y_pred,y_test))
```

```
Training Accuracy Score : 0.7818181818181819
```

```
Training Confusion Matrix  : [[20  5]
 [ 7 23]]
```